



BERN UNIVERSITY OF APPLIED SCIENCES (BFH)

BACHELOR THESIS

CHVote Demonstrator

Authored by Kevin HÄNI <kevin.haeni@gmail.com>
Yannick DENZER <yannick@denzer.ch>
Supervised by Prof. Dr. Rolf HAENNI <rolf.haenni@bfh.ch>
Prof. Dr. Philipp LOCHER <philipp.locher@bfh.ch>

Bern, December 27, 2017

Contents

1	Introduction	7
1.1	Electronic voting	7
1.2	CHVote	8
1.3	Project task	8
2	Project Management	11
2.1	Goals	11
2.1.1	General Requirements	11
2.1.2	Election-Overview	12
2.1.3	Election Administrator	12
2.1.4	Printing Authority	12
2.1.5	Election Authority	12
2.1.6	Voter	13
2.1.7	Bulletin Board	13
2.1.8	Out-of-Scope	13
2.2	Time Schedule & Implementation Phases	14
2.3	Use Cases	16
3	CHVote Protocol	17
3.1	Actors	17
3.2	Pre-Election & Voting Cards	18
3.3	Vote casting with oblivious transfers	18
3.4	Anonymity with a e-ncryption Mix-Net	19
4	Product Description	21
4.1	Application Overview	21
4.2	Design	25
5	Technical Implementation	27
5.1	Technology & Language Decisions	28
5.2	Architecture	29
5.3	Back-End	30
5.3.1	VoteService	31
5.3.2	Data-Sync Service	35
5.3.3	REST API	37
5.4	Crypto-Library	39
5.4.1	File structure	39
5.4.2	Public Parameters	40

5.4.3	Coding Style	40
5.4.4	Return Types	42
5.5	Front-End	43
5.5.1	Components	43
5.5.2	Centralized Data-Store & Flux Pattern	46
5.5.3	Internationalization (i18n)	46
5.5.4	Development Environment	47
5.6	Challenges	47
5.6.1	Websocket subscription concept	47
6	Conclusion	49
6.1	Python issues	49
6.2	Reflection	50
7	Appendix	53
7.1	Use Cases	53

Management Summary

1 Introduction

With the ongoing digital transformation it seems just a matter of time until the paper based ballots will be replaced with electronic voting (e-voting). Recent polls show that the majority of the swiss population is interested in having the possibility to vote online ¹. However, e-voting is a very complex topic and designing a secure e-voting protocol is known to be notoriously challenging in terms of IT-security and cryptography.

1.1 Electronic voting

First trials with e-voting in Switzerland date back to 2003. Since 2015, it is possible for Swiss citizens registered in the cantons Geneva and Neuchâtel and living abroad, to vote electronically. These systems were available only for a limited electorate size since they did not yet meet the requirements in terms of security and transparency, to be accepted as a secure E-Voting platform on a nationwide scale.

An e-voting system must satisfy a large variety of security requirements set up by the government, including:

- **Fairness:** It is not possible for any person (including the participating parties of an e-voting system) to learn the intermediate result or outcome of an election before the result has been officially tallied and published to a public board.
- **Privacy:** No one can find out information about a voters selection. This implies that a voters ballot must be encrypted before it leaves the voters client and until the election is tallied.
- **Authenticity:** All voters must be authenticated as eligible voters in order to cast a vote
- **Soundness:** Only valid votes are being tallied. If a voter selects more candidates than he is allowed or less than he is supposed to, the vote must not be counted.
- **Robustness:** An e-voting system detects cheating actors.

There exist many different concepts and e-voting protocols that cover most of these requirements. However, most of the existing solutions behave as some kind of black-box and keep the

¹https://www.swissinfo.ch/ger/umfrage_grosse-zustimmung-zu-e-voting-trotz-sicherheitsbedenken/42457426

internals hidden, such that a voter cannot be sure, whether his vote has been recorded correctly and counted in the final result.

One of the requirements that is hardest to achieve is the **universal verification**: A good e-voting system must be transparent and allow an external person to verify, that every protocol participant has abided by the protocol and that all and only valid votes have been counted correctly.

Another common problem is the **insecure platform problem**: If a voters computer is affected by malware, the vote casting process is no longer under the voters control and the candidate selection could be possibly manipulated without the voters notice. It must be **individually verifiable** to every voter that his intended vote has been recorded, while at the same time, the voters privacy must still be ensured at all times.

If an e-voting system shall be available to more than 30% or 50% of the respective cantonal electorate, it must be both individually as well as universally verifiable.

1.2 CHVote

Some of the mentioned requirements might sound as a paradoxon at first. However, they can be solved by advanced cryptography. A contract was formed between the state of Geneva and the the Institute for Security in the Information Society (RISIS) of the Bern University of Applied Sciences to work out a new protocol which does meet the complex requirements set up by the swiss government. In 2017, Rolf Haenni, Philipp Locher and Reto E. Koenig published the resulting specification document and a proof-of-concept has been successfully implemented by the State of Geneva. Geneva is currently developing a new version of their e-voting system with the name "CHVote", based on this specification. Other cantons, currently St. Gallen, Aargau, Bern and Lucerne have announced their interest in cooperating with Geneva and might use their system in future.

The CHVote specifications document is publicly available and describes not only the theoretical background and the algorithms, but also provides pseudo-code for the approximately 75 algorithms that are needed to implement their protocol.

1.3 Project task

An additional problem e-voting is confronted, is that understanding such a complex protocol isn't easy without good knowledge in cryptography and mathematics. This might be one of the reasons why many people still do not trust electronic voting systems.

In close consultation with the authors of the CHVote specification, we were looking for an interesting project for our bachelor thesis in the thematic field of e-voting. We decided to study the CHVote specification, implement the protocol and build an application that allows to get

a hands-on experience with the Geneva e-voting system and that could be used to show and explain to an audience how the future of voting in Switzerland might possibly look like.

Bachelorthesis-Aufgabe

Implementierung des Genfer E-Voting Systems

ID	IHNR1-1-17
Studierende	Kevin Häni Yannick Denzer
Betreuer	Dr. Rolf Haenni Dr. Philipp Locher
Experten	Han van der Kleij
Aufgabe	<p>Der Kanton Genf bietet seit vielen Jahren ein E-Voting System an, mit welchem die Bürgerinnen und Bürger ihre Stimme elektronisch abgeben können. Das zur Zeit eingesetzte System genügt aber den aktuellen Sicherheitsanforderungen des Bundes nicht mehr. Aus diesem Grund wurde vor zwei Jahren die Entwicklung eines neuen Systems beschlossen, in welchem das Wahlergebnis von unabhängigen Stellen überprüft werden kann.</p> <p>Eine Forschungsgruppe der Berner Fachhochschule wurde beauftragt, für dieses neue System eine genaue kryptographische Spezifikation zu entwickeln. Aufgrund dieser wird zur Zeit in Genf das neue System entwickelt. Da die Spezifikation ein öffentliches Dokument ist, lässt sich das System exakt nachbauen. Das Bauen eines solchen Replikats des neuen Genfer E-Voting Systems ist das Ziel dieser Bachelorarbeit. Das erwartete Resultat soll dann unter anderem zu Test- und Demo-Zwecken eingesetzt werden, um der Öffentlichkeit ein besseres Verständnis des Genfers System zu vermitteln.</p>

Figure 1.1: Bachelor thesis task

Our project task required us to read and get a good understanding of the CHVote specification. As preparatory work during the course "Project 2", which we have finished just before the start of our bachelor thesis, we have implemented the approximately 65 algorithms which were defined as pseudo-code in the CHVote specification document. The resulting set of algorithms has been used as a library for implementing the protocol on which our application will be built on.

Based on this rather general task description, there were several possibilities regarding the final product, such as a realistic prototype, a verifier software for the implementation which is being developed in Geneva, or a demonstrator-application that targets the educational problem.

Ultimately we have decided for the last options: Our goal was to develop a web-based application that allows to visualize every step of a CHVote election event, from the pre-election tasks like generating the electorate data, to casting and confirming ballots from a voters point

of view, to the post-election processes like mixing, decryption and tallying. Opposed to a realistic prototype, the focus on our project is not to implement an e-voting-system that is totally realistic and secure from an architectural perspective, but more on the visualization.

The authors of the specification have the intention of using our application for future demonstrations of their protocol to create a better understanding and acceptance for e-voting.

In chapter 2, we will further discuss the goals and requirements that we have set for this project. Additionally, it covers the time planing and project methodology of our project.

Chapter 3 will describe the final product and the concepts behind our application

Chapter 4 covers the technical aspects regarding the implementation, such as the architecture, the language and technology decisions. Later sections contain detailed information about the internals of each component of our application and the challenges we were facing during the implementation phase.

Chapter 5 will close this document with a short summary and a reflection about the project.

2 Project Management

In this chapter we are going to describe several aspects regarding the planing of our project, such as our project methods, the requirements and use cases as well as a short concept of our application.

Since this wasn't a project where the project scope and goals were strictly defined and clear from the beginning, it was very important to elaborate the requirements in close collaboration with our supervisors early on. This is why we decided our project would be very suitable for an agile kind of project method. Because of the small team, consisting of only two members, we didn't chose a particular project method like SCRUM, as this would have probably caused more overhead than actual benefits.

We have therefore set up regular meetings (usually once every 2 weeks) with our supervisors to discuss our ideas and get feedback about the current progress.

2.1 Goals

As the first step, we have elaborated the goals and requirements for our application. We have structured the requirements into groups that correspond to the actors of the CHVote protocol and therefore the views our application is going to consist of.

Some requirements affect multiple or all actors and are therefore listed as "General Requirements". We have also added a priority and requirement type to simplify the planing and time management.

2.1.1 General Requirements

	Description	Type	Prio.	Phase	Status
R1	The CHVote protocol is implemented as specified in the latest specification document. The only exclusion are the algorithms for channel security.	Must	High	1	Done
R2	The application is web-based shows updates within the same demo-election in real-time.	Must	High	1	Done
R3	The system supports 1-out-of-3 type of elections (e.g. elect 1 of 3 possible candidates)	Must	High	1	Done
R4	The system supports multiple parallel elections	Must	High	1	

R5	Users can create new elections	Must	High	1	Done
R6	The system supports internationalization. Providing more than one language is not required.	Must	Med.	1	
R7	The system can handle k-out-of-n type of elections	Can	Med.	1	

2.1.2 Election-Overview

	Description	Type	Prio.	Phase	Status
R8	The overview shows which phase the election is currently in	Must	High	2	
R9	A graphical scheme of the chVote protocol gives an overview of all participating parties	Must	Med.	2	

2.1.3 Election Administrator

	Description	Type	Prio.	Phase	Status
R10	An election can be set up by providing all required information such as the candidates, number of parallel voters, the number of voters and the number of selections (simplified JSON input)	Must	High	1	Done
R11	The election can be set up without entering the parameters in JSON format and allows easier set up of elections with multiple parallel election events	Must	Low	2	
R12	The election administrator view allows to perform the tallying and displays the final result of an election in numbers and a pie chart	Must	High	1	
R13	During election setup, the security parameters can be chosen from a set of predefined parameters	Can	Low	2	

2.1.4 Printing Authority

	Description	Type	Prio.	Phase	Status
R14	Users can generate and display voting cards for an election.	Must	High	1	Done
R15	Voting cards hide sensitive information behind a scratch card	Can	Med.	2	

2.1.5 Election Authority

	Description	Type	Prio.	Phase	Status
R16	The election authority view shows all information known to an election authority	Must	High	1	Done

R17	After a voter has submitted a ballot, all election authorities can check and respond to the voters submission	Must	High	1	Done
R18	In the post-election phase, all election authorities can perform the mixing and decryption tasks	Must	High	2	
R19	Each authority can optionally processes all tasks automatically	Can	High	2	Partially done

2.1.6 Voter

	Description	Type	Prio.	Phase	Status
R20	Users are able to go through the whole vote-casting process for every voter	Must	High	1	Done
R21	The voting card of a voter is displayed on screen. The voting and confirmation codes can be copied into the input textfields by double clicking	Must	Med.	1	

2.1.7 Bulletin Board

	Description	Type	Prio.	Phase	Status
R22	The bulletin board view shows what information is publicly available	Must	High	1	Done
R23	The bulletin board view is extended with verification-functionality	Can	Low	2	

2.1.8 Out-of-Scope

The following topics are considered out-of-scope for the duration of our project:

- The goal of our project isn't to build a realistic prototype. Therefore, the whole back-end will run on a single server while in reality, there would be components running on distributed servers. Another difference between our implementation and a real implementation is that we generate the ballots on the server. In reality, the ballots would have to be generated on the client for security reasons. This however would require us to rewrite many of the already implemented algorithms in JavaScript.
- The protocol takes into account that not all voters might be eligible to vote in all elections of a given election event (eligibility matrix). For simplicity, we assume that all voters are eligible to vote in every election for our application.
- Message level encryption and signature based integrity protection are very important in a real implementation of a evoting-system and are also described in the CHVote specification. However, as our system is only used for demonstration purposes and as we do not

have a distributed infrastructure, there is no real need for channel security in our project.

- Providing more one language is also out-of-scope. If there is enough time, a second language might be provided optionally.

2.2 Time Schedule & Implementation Phases

As the next step we created a time schedule and structured our project into smaller units.

The actual implementation phase has been broken down into two phases:

- Phase 1 involves the implementation of all high priority must-requirements, so, basically bringing the application into a state that allows to visualize the whole CHVote election process. We have agreed that after phase 1, some areas of the user interface would still be in a rather primitive state (eg. user inputs are not validated and require a more technical type of input).
- For phase 2 we planned on implementing the can-requirements and the must-requirements with lower priority.

This approach reduced the risk of technical limitations of our architecture to remain hidden until later that would have resulted in time consuming architectural changes.

Project Tasks			38	39	40	41	42	43	44	45	46	47	48	49	50	51	52	53	54
Documentation																			
Update documentation and journal	soll																		
	ist																		
Concept																			
Working out goals / requirements	soll																		
	ist																		
Concept	soll																		
	ist																		
Implementation																			
Update CHVote cryptolibrary to the latest specification	soll																		
	ist																		
Prototyping (proof of concept)	soll																		
	ist																		
Iteration 1																			
Implement backend API	soll																		
	ist																		
Implement frontend: Pre-election	soll																		
	ist																		
Implement frontend: Election	soll																		
	ist																		
Implement frontend: Post-election	soll																		
	ist																		
Iteration 2																			
Automation of election authority (R19)	soll																		
	ist																		
Voting Card layout, Scratchcard (R15)	soll																		
	ist																		
Election Overview (R8, R9)	soll																		
	ist																		
k-out-of-n election types (R7)	soll																		
	ist																		
Improving the input forms and layout (R11)	soll																		
	ist																		
Flexible security level (R13)	soll																		
	ist																		
Verifier functionality (R23) and reserve	soll																		
	ist																		
Project completion																			
Poster, Final Day, Presentation	soll																		
	ist																		

Table 2.8: Project schedule

From our requirements we have defined the following mile-stones:

- M1: Finishing the implementation of the CHVote cryptolibrary
- M2: Upon successful creation of a proof-of-concept / prototype for our application
- M3: After finishing implementation phase 1
- M4: Finishing implementation phase 2 including the can-requirements
- M5: After finishing our documentation

2.3 Use Cases

The next step was to create use cases. As an example we show one of the use cases, for the complete list of use cases we refer to the appendix.

Table 2.9: Use Case «Casting of a vote»

Use Case	Casting of a vote
Primary Actor	Voter
Description	The voter can cast a vote by selecting his favored candidate(s) and his voting code
Precondition	<ul style="list-style-type: none"> • The election has the status "Election Phase" • A voter is selected in the "Voter"-view • The voter has the status "Vote Casting Phase"
Postcondition	The first election authority receives a "Ballot-check task"
Main path (M)	<ol style="list-style-type: none"> 1. The voter visits the "Voter"-view and select his voter object 2. The system demands a selection of the candidates and the voter's voting code 3. The voter clicks on "Cast ballot"

3 CHVote Protocol

The protocol our project is based on, does not originate from us! The concept and specifications have been created by Rolf Haenni, Philipp Locher and Reto E. Koenig of the Institute for Security in the Information Society (RISIS) of the Bern University of Applied Sciences. For this project, we have implemented the protocol according to their specification. In this section we summarize the most important aspects of the protocol and establish the terminology for better understanding this document and our application

3.1 Actors

A **voter** is a person who is eligible to vote in his respective state. Every voter is assigned a **counting circle** which typically corresponds to the voters municipal and is required for statistical purposes. For authentication purposes, every voter must possess a voting card that has been sent to him prior to an election event that contains codes used to identify the voter during the vote casting process.

The **Election Administrator**, typically a person of the government, is responsible for setting up the election event by providing the required information such as the candidates or the voters and initiates the generation of the cryptographic electorate data. He is also responsible for tallying the election and publishing the final results.

The **election authorities** can be seen as some kind of independent election observers. In a CHVote election event there are always multiple election authorities in order to avoid having to trust a single authority. The authorities are involved in almost every step of the protocol, starting from the generation of their shares of the electorate data, checking and responding to new ballots, as well as in the mixing (a measure to ensure anonymity) and decryption phases. The public key that is used for encrypting the ballots has been jointly generated by all election authorities. Therefore in order to decrypt the ballots, all authorities must work together and provide their share of the private key.

The measure of multiple authorities participating in the whole e-voting process establishes a **distribution of trust** and ensures security of the whole election process as long as at least one election authority can be trusted.

The **Bulletin Board** acts as a central board over which most communication is done and where all the public data is stored. By publishing all data that is not secret onto a public board, including the list of encrypted ballots, and functionality to verify the correctness of

these data, the protocol is making a big step towards the demanded transparency and the universal verification.

The **Printing Authority** is responsible for printing the voting cards for all voters. Since the printing authority needs to be in possession of all the secret voting codes, it is a very sensitive point in the protocol. The printed voting cards are handled over to a trusted mailing service for delivery.

3.2 Pre-Election & Voting Cards

Before the actual election phase, the vote administrator sets up the election event by entering all required parameters, including the voters, the elections with their corresponding candidates and the number of candidates a voter can select in every election. All election authorities jointly generate the cryptographic data for the whole electorate from which the voting card information is derived. The printing authority combines the information from all election authorities, prints the voting cards and delivers them to the voters by a trusted mailing service.

A voting card contains several codes, namely:

- a voting code
- a confirmation code
- a finalization code
- one verification code for every candidate

The **voting and confirmation codes** are authentication codes and are used to authenticate the voter twice during the vote casting process; the first time with the voting code when he casts his vote, and a second time for confirming his vote. A second authentication code is required because otherwise an attacker who infected a voters computer with malware could just confirm a vote on behalf of the voter after he manipulated the candidate selection and could therefore skip the whole verification process.

3.3 Vote casting with oblivious transfers

As mentioned earlier, one of the big challenges of an evoting-protocol is how it deals with the insecure platform problem: A voting platform that is infected with malware poses the risk that an attacker can manipulate the candidate selection on the vote-casting page after the voter has entered his voting code.

The CHVote specification suggests a "Cast-as-intended verification"-step to allow voters to detect this kind of manipulation: When a voter enters his voting code and the indices of his

avored candidates, the voting client forms a **ballot** containing the voters selection encrypted with the authorities public key, his public voting credential derived from the voting code, and a **non-interactive zero knowledge proof** which proofs that the voter has formed the ballot according to the protocol and that he has been in knowledge of his voting code, without revealing any information about the voting code.

Every election authority has to check the voters public voting credential, the validity of the ballot proof and that the voter hasn't already cast a vote. The encrypted selection also serves as a query for an oblivious transfer. A **k-out-of-n oblivious transfer** allows a client to query a server for k messages, without the server knowing what messages the client requested, and without the client learning anything about the other $n - k$ messages. Adapted to the CHVote protocol: The voting client queries the authorities for the corresponding verification codes of the selected candidates, without the authorities learning which candidates the voter has selected, and without revealing any information about the other candidates.

The voter then checks if the returned verification codes match the codes of the candidates he has chosen on the printed voting card. If the selection was somehow manipulated by malware, the returned verification codes would not match the printed ones and the voter would have to abort the vote casting process. This way the integrity of the vote casting process can be guaranteed even in the presence of malware. In such a case, privacy on the other hand cannot be protected since the malware will learn the plaintext of the voters selection.

Another feature the protocol supports, is that an election event can consist of t multiple parallel elections. In such cases, the voter has to submit a single ballot, which contains his candidate selection for all parallel elections. This raises the question, how the system can verify that a voter has chosen exactly the correct number of candidate in each election, and not for example one less in the first, and one additional candidate in the second election.

The specification suggests the following trick: Assuming an election consists of two parallel elections ($t = 2$) with 3 candidates each ($n_1 = 3, n_2 = 3$), of which a voter can select one candidate in each election ($k_1 = k_2 = 1$). The verification codes are derived from $n = \sum_{j=1}^t n_j$ random points on t polynomials (one for every election event j) of degree $k_j - 1$, that each election authority has chosen randomly prior to the election. By learning exactly $k = \sum_{n=1}^t k_j$ points on these polynomials, the voting client is able to reproduce these polynomials and therefore is able to calculate a particular point with $x = 0$ on these polynomials. The corresponding y values are incorporated into the second credential from which the confirmation code is derived. As a result, only if a voter has been able to reconstruct these polynomials with the returned points by submitting a valid candidate selection, he will be able to confirm the vote that he casted.

3.4 Anonymity with a e-ncryption Mix-Net

As mentioned earlier, both the election authorities as well as the bulletin board keep track of all the casted ballots, which contains the voters candidate selection encrypted under the public key that has been jointly generated by all election authorities. Up to this point, the ballots are

still linked to the voters, since the voter identifier is required for the confirmation process and to avoid that a voter can cast multiple ballots. If the ballots were decrypted now, this would conflict with the anonymity and the privacy of the voters.

As the first step of the post-election phase, the first election authority extracts the encrypted candidate selections from the ballots. In order to anonymize this list of encryptions, the protocol suggests a "mixing-process", in which every authority performs a cryptographic shuffle with a random permutation. Additionally, all the encryptions are re-encrypted such that they receive a new ciphertext. This re-encryption is done by using the multiplicative homomorphic property of the ElGamal encryption scheme that is used for encrypting the ballots. A **multiplicative homomorphic encryption** scheme allows to perform multiplications on the ciphertext such that:

$$Enc(a) \cdot Enc(b) = Enc(a \cdot b)$$

The specification suggests multiplying the encrypted selection with the encryption of the neutral element 1 since this yields a new ciphertext for the exact same plaintext.

$$Enc(a) \cdot Enc(1) = Enc(a)$$

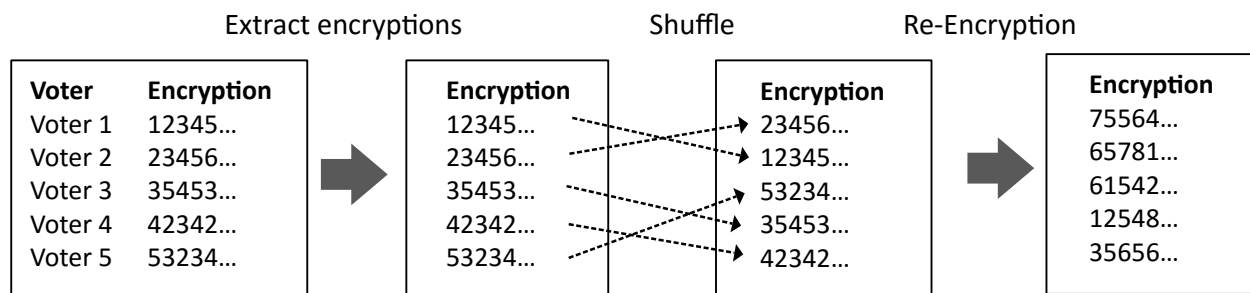


Figure 3.1: During the mixing phase, the encryptions are extracted from the list of ballots and then shuffled with a random permutation and re-encrypted by all election authorities. This measure ensures the anonymity of the votes.

While the extraction only needs to be done by the first election authority, every election authority is sequentially performing the shuffle and re-encryption on the mixed list of the previous election authority.

To prevent election authorities from cheating and not performing the mixing as specified, a proof (according to Wikström's shuffle proof) is calculated which will be verified by all election authorities before decryption. If one of these **shuffle-proofs** fails, the election process can not proceed.

4 Product Description

In this chapter we describe the product of our bachelor thesis and the main component of our application - the demonstrator web-application (front-end). We will focus mainly on the non-technical concepts as the next chapter will explain the technical implementation of the product.

4.1 Application Overview

In an essence, all the functionality of our application revolves around visualizing an election event of the CHVote protocol and guiding the users through the several phases of such an election event. A typical CHVote election event can be broken down into phases shown in figure 4.1

To allow the users of our application to experience and get insights about the workings of the protocol, we wanted let them take the perspective of every actor of the protocol at any given time and during every phase of the election event. This is why we have divided our application into separate **views**, one for every actor:

Every actor involved in a CHVote election event has it's own set of data to be displayed and specific tasks and use-cases it has to perform. Most use-cases are only available during a particular phase.

- **Election Overview:** The election overview shows in what phase the chosen election event is currently in. Additionally, a graphical schema shows how all the actors of the e-voting ecosystem are connected and who is currently in charge of performing an action.
- **Election Administrator View:** The view of the election administrator allows a user to set-up an election event, by configuring the number of voters, the elections including the candidates and the number of selections per election. Instead of providing these information every time an election event is set up, previously defined election presets can be applied or the parameters can be generated randomly.

The election administrator view is also the place where the elections can be tallied and the final result is determined during the post-election phase.

- **Printing Authority View:** In the printing authority view, the voting cards can be printed and displayed for every voter. Additionally, the voting cards can be sent to




Figure 4.1: In our application, an election event consists of 7 different phases (green). During the election phase, a voter can be in 3 different phases (blue)




Figure 4.2: The election overview guides the user through the several phases of an election event and indicates what phases have been finished and what the next steps are.

the voters.

- **Election Authority View:** The election authority view first lets a user choose one of the three election authorities he wants to observe. On the top, new tasks will pop up whenever an election authority needs to perform a specific task such as ballot- or confirmation checking or mixing and decryption.

Additionally, the view shows the data that an election authority knows. This concept of diving a view into tasks and data, can best been seen in figure 4.3 and has been used throughout the application and in almost every view.

- **Voter View** In the voter view, the vote casting process can be done for every voter that has been previously generated. The view displays a candidate selection and code input form in the left, and a voters voting card in the right side. The voting cards codes are hidden behind a scratchcard if they are sensitive and private to the voter and can be applied to the form by clicking on them after they have been scratched open.

The voter view additionally has to consider the current phase of a voter e.g. if he has already casted a vote, if he already confirmed his vote or if he aborted the vote casting process.

- **Bulletin Board View** The bulletin board view always shows all the data that has been appended to the bulletin board by the other actors, such as the pre-election data, the ballots that the voters have casted, as well as all the proofs generated during the post-election phase.
- **Verifier View:** The view of the verifier becomes visible after the election result has been published to the bulletin board and the election event has reached its final stage. By clicking on the verify button, several checks can be executed and the result is display on the page.

All the views are accessible from a tab view displayed on the top of the page which serves as our standard mean for navigation, see figure 4.4.

From the use-cases, we tried to figure out all commonalities between the views: The views typically display the information known by the respective actor. Especially the bulletin board and election authorities will contain lots of information to be displayed. Most of the views have distinct tasks to be executed by the respective actor, such as casting a ballot in the voters view or confirming ballots from an election authorities view.

The content that a view displays, or in other words, the functionality an actor has access to, depend on the phase the election event is currently in: During the pre-election phase, the vote administrator needs to able to set up the election while in the tallying phase, he must be able to tally and determine the final result.

The following matrix shows all the possible combinations of views and phases.

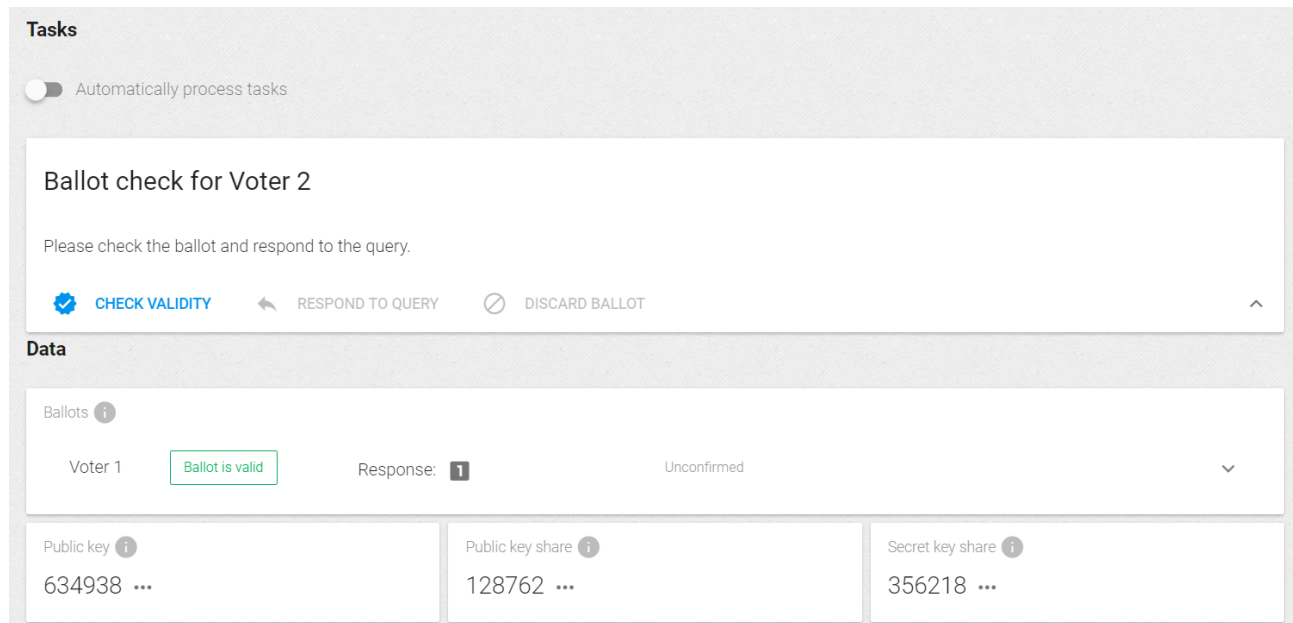


Figure 4.3: The election authority view as an example of how we generally structure our pages: Most views are divided into a tasks part, which allows to perform the tasks of an actor, and a data part which shows what information about an election event the participant is in possession of.




Figure 4.4: The tab beyond the page header is the main control for navigating through the different pages. It also indicates where some interaction is required next by displaying an exclamation-mark icon in the respective actors tab

4.2 Design

Given the rather large amount of complex data to be displayed, the main challenge of the project has been a well designed user interface that allows to display all important information while maintaining a clear overview.

To achieve this goal we tried to keep our design very minimalistic and follow the Google material design guidelines as much as possible by choosing an appropriate UI component framework. Even though mobile compatibility hasn't been a requirement, still it was nevertheless our goal to make the layout as responsive as possible such that it can at least be used from a mobile phone. The web-application is definitely optimized for desktop use.

Throughout our application we tried to establish common concepts regarding the look and feel and on how to display our data. One popular layout-concept of the google material guidelines is the card layout. Cards can be easily integrated in a responsive grid system, look modern and allow to visually group data. In addition, we used pushover menus, tooltips and popups as they made it possible to hide lesser relevant information by default and display it only on demand of the user.

Before we started implementing our application, we have created mockups for most of the views, to discuss our ideas with our supervisors and incorporated their feed as well. The following screenshots are an extract of the mockups in which we tried to visualize how we imagined the resulting application to look like during the conceptual phase.

Our conceptual work also involved the development of a small prototype / proof of concept, in which we have implemented one usecase in a reduced extent with the envisaged frameworks and technologies to evaluate the technical feasibility. For more information about the languages and frameworks we have decided on, we refer to the next chapter.

5 Technical Implementation

In this chapter, we will describe how we implemented our application. We will start explaining the architecture from a high-level perspective with each component being a black-box. Later sections of this chapter will then further describe the internals of every component.

We have implemented our application following a "single page application (SPA)" architecture. Some reasons that lead to this decision:

- Due to our personal interest in JavaScript and our intention of improving our knowledge in this language, we wanted a significant part of the development to be done client-side
- We imagined the state handling to be easier with a SPA than having to pass around cookies and session data between every request
- A SPA seemed great for building modern looking, intuitive and responsive user interfaces

From a high level perspective and following the SPA architecture, we have structured our application into 3 main components, as shown in figure 5.1:




Figure 5.1: From a high level perspective, ur application consists of 3 main components: The front-end (web-application), the back-end, and the crypto-library.

- The **front-end** / web-application is where all the functionality of our back-end is consumed and where a typical CHVote election event is visualized for the users. It is therefore the most important component of our application. The backend is developed in response to what is needed by the frontend.

- The **back-end** consists of several components that make use of the CHVote crypto-library to build an actual e-voting ecosystem and providing an API to manipulate the data as well as a data-synchronization service to push data from the back-end to the web-clients. All the sub-components of the backend run on a single server.
- The **CHVote crypto-library** is the result of our "Project 2" course project which we have finished before our bachelor thesis. This library contains all the algorithms that are specified as pseudocode in the specification document.

5.1 Technology & Language Decisions

When we implemented the CHVote crypto-library, we have evaluated and decided to use Python. Since Java has already been used by the team in Geneva, we wanted to use a different language as it was also desirable to prove that the CHVote specification can be implemented regardless the programming language. Python seemed like a rather suitable language for our project due to the following reasons:

- python is a mature language with lots of libraries
- python allows programs to be written in a compact and readable style, for example by supporting tuples
- as the protocol wasn't completely specified at that point and has still been undergoing some changes, we wanted to use a language in which we can adapt changes quickly and easily
- native support for large integers (*BigInts*) and bindings for the GMP¹ library
- supports a lot of platforms
- many popular web development frameworks are available

Throughout the project not all of the reason above turned out to be true or ideal. The drawbacks that we have experienced during the implementation of this project will be discussed in the last chapter of this document.

Since we used the crypto-library in our back-end, Python was also the obvious language for the whole back-end. Python offers a wide variety of frameworks for building web-services. Since we planed on building a single-page-application for the client, we chose the lightweight micro-framework flask for building a restful web-service and the data synchronization service.

For our **front-end** (web application) we evaluated several single-page application frameworks. VueJS is a new, modern and lightweight SPA framework that in contrast to Angular has a much flatter learning curve but still offered all the functionality that we needed. The VueJS add-on

¹GMP is a free library for arbitrary precision arithmetic, operating on signed integers, rational numbers, and floating-point numbers, see <https://gmplib.org/>.

Vuex enabled us to establish a data-store pattern in our front-end, which makes it possible to have a copy of the back-end data-store in our web application which is synchronized in real-time through web-sockets.

Socket.io simplifies the usage of web-sockets and offers fallback technologies such as long-polling in case web-socket is not supported by either the browser or the web-server. Both Flask as well as VueJS have plug-ins that support and integrate socket.io.

For persisting the state of an election, we decided to use mongoDB. The reason behind this choice will be described in more details later on.

5.2 Architecture

The core of our application is the VoteService component in our backend which implements the e-voting protocol by utilizing all algorithms of the CHVote crypto-library according to the CHVote specification. The VoteService component internally holds the state of a whole CHVote election event and exposes functions to manipulate this state at a granularity required by our web application to implement all use cases. For example: The VoteService contains a list of ballots and exposes functions to cast a new ballot, which will generate a new ballot according to the protocol, by calling the CHVote crypto-library, and adds the ballot to the list.

On top of the VoteService we have implemented a REST service that acts a facade to the VoteService component and makes its functionality available as an API to our web-clients. The REST service also has to initialize the VoteService by loading and persisting it's state from and to the database between each API call, since each API call is stateless.

One of the requirements is that all clients must be informed in real-time about mutations of the election state made by other participants. To achieve this, we have implemented a data sync service which allows to push the state of an election event to the web-clients by using the websocket protocol. This service is triggered by the REST service after every API call to push the delta between the old and the new state to the clients.

To establish a proper separation of concerns, the state of the VoteService is always sent to the client via the data sync service. The REST API only returns success or error codes or information that is required in response to some particular API call, and never state objects. On the other hand, the data sync service does never manipulate the state of the VoteService and is solely responsible for sending s to the client.

From the clients point of view, the web-client contains a copy of the whole VoteService state in a local data-store. This store is initially populated when the web application is initialized with an election. Whenever the state of the election event / VoteService changes, the data sync service pushes the new data to the web-client. A local mutation handler within the web application handles those messages and writes the new data into the local data-store.

Since the components that the pages of the web application are built of, are bound to the

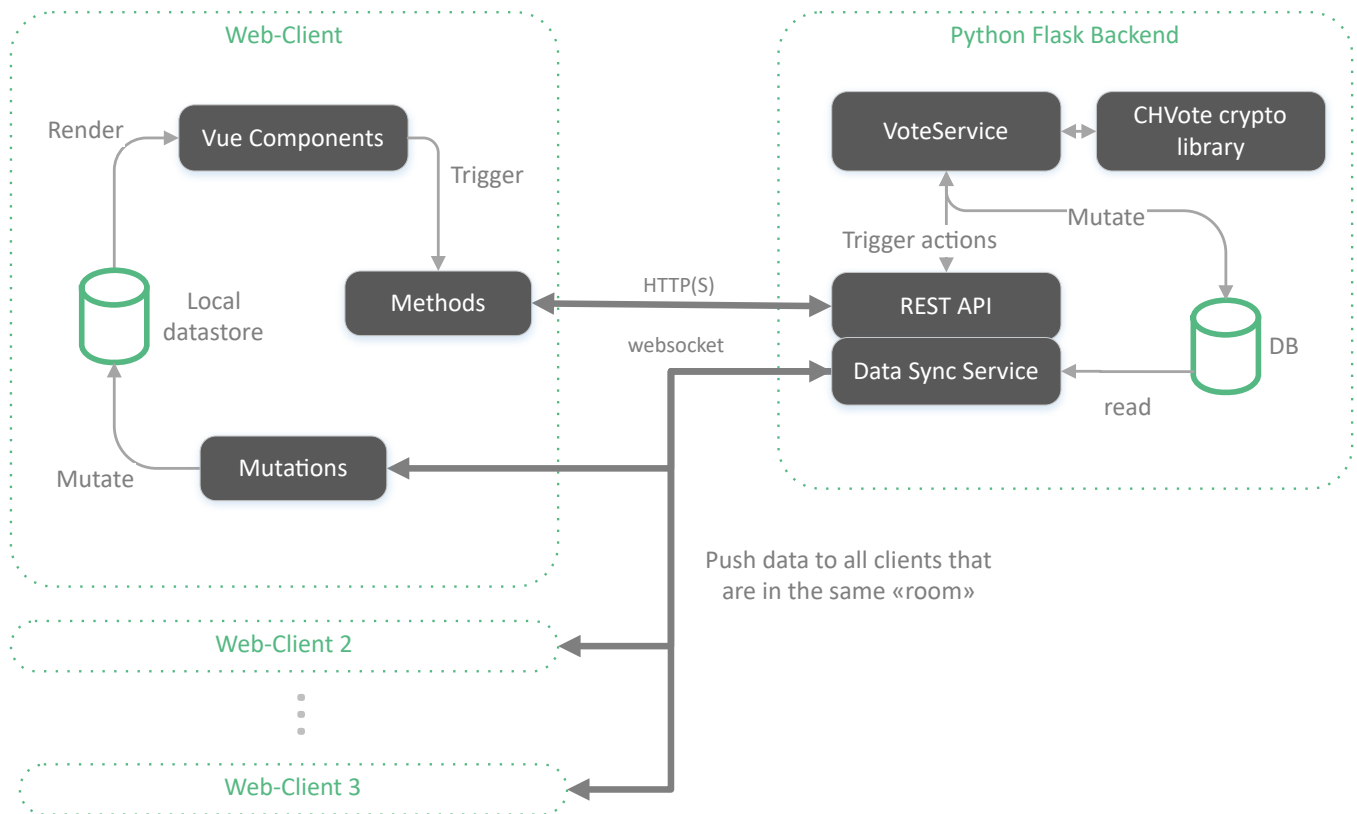


Figure 5.2: Our architecture with the front- and the back-end. Both sides contain a data store or data base that stores the current state of an election event and that is directly bound to the user controls. Whenever some action is called on the server side by calling our REST API, the changes are written to the database and synchronized to all clients over web-sockets.

local data-store, all mutations are automatically reflected to the user. From those pages, the REST API can again be called to trigger some election specific action on the back end. The resulting state change is again being pushed to all clients. The responsible client that performed the API request will additionally receive a success code, or an error message in case of an error over HTTPS.

The architecture with all corresponding components is also shown in figure 5.2

5.3 Back-End

In this section we describe the internals of the back-end services.

5.3.1 VoteService

The VoteService builds a simplified e-voting ecosystem that provides functionality to conduct an election event by internally representing the state of a whole election event. In reality, an e-voting system based on the CHVote specification would consist of multiple separate applications/services for every participant, such as a bulletin-board service for appending data to the public board, or services for every election authority which performing their tasks. Additionally, several steps of the protocol would have to be executed on the client-side, such as the generation of a voters ballots. In our application, it was reasonable to have the protocol functionality of all these parties combined within a single service.

We have decided on this VoteService centric approach mainly because we wanted to keep the whole protocol implementation within a central component and wanted to avoid having protocol logic both in the client as well as in the back-end. As an advantage, if the protocol is undergoing any changes or if our demonstrator application should be extended to support a different e-voting protocol in future, only the VoteService component (and of course its dependencies as the crypto-library) will be affected or must be replaced. It also allowed us to implement the CHVote protocol almost identically as described in the specification. Because we had access to every actors state and functions from within our VoteService object and could pass data from one actor to another as simple as setting an object property to some value.

Internally, we have divided the VoteService into actors and states, as shown in figure VoteService: One actor for every protocol participant, providing the functionality a participant is responsible for, and a corresponding state-object for every actor, representing the participants current state within a given election event.

The distinction between the actors and their states allowed us to easily serialize the state objects to JSON (a format that can be easily interpreted by the front-end) and made it easy to load and persist the state from and to the database. This measure was also necessary because of the way how we implemented our data synchronization between the clients and the back-end: By comparing and determining the differences between the state object before and after some VoteService actions, we can automatically find out the mutations that has been done to the state objects using the JSON Patch standard and generate operations to patch the clients local datastore in the same way. More about this technology can be found in the data-sync section.

The only common functionality between every state object, is the ability to serialize the object to a JSON string. For this reason we had to write a custom transformer which tells the JSON parser how to serialize data-types such as mpz, bytearrays and custom classes. Luckily, python offers a way to easily serialize any custom object. By calling `object.__dict__` we can convert an object into a dictionary, as long as the transformer is able to serialize all properties of the object.

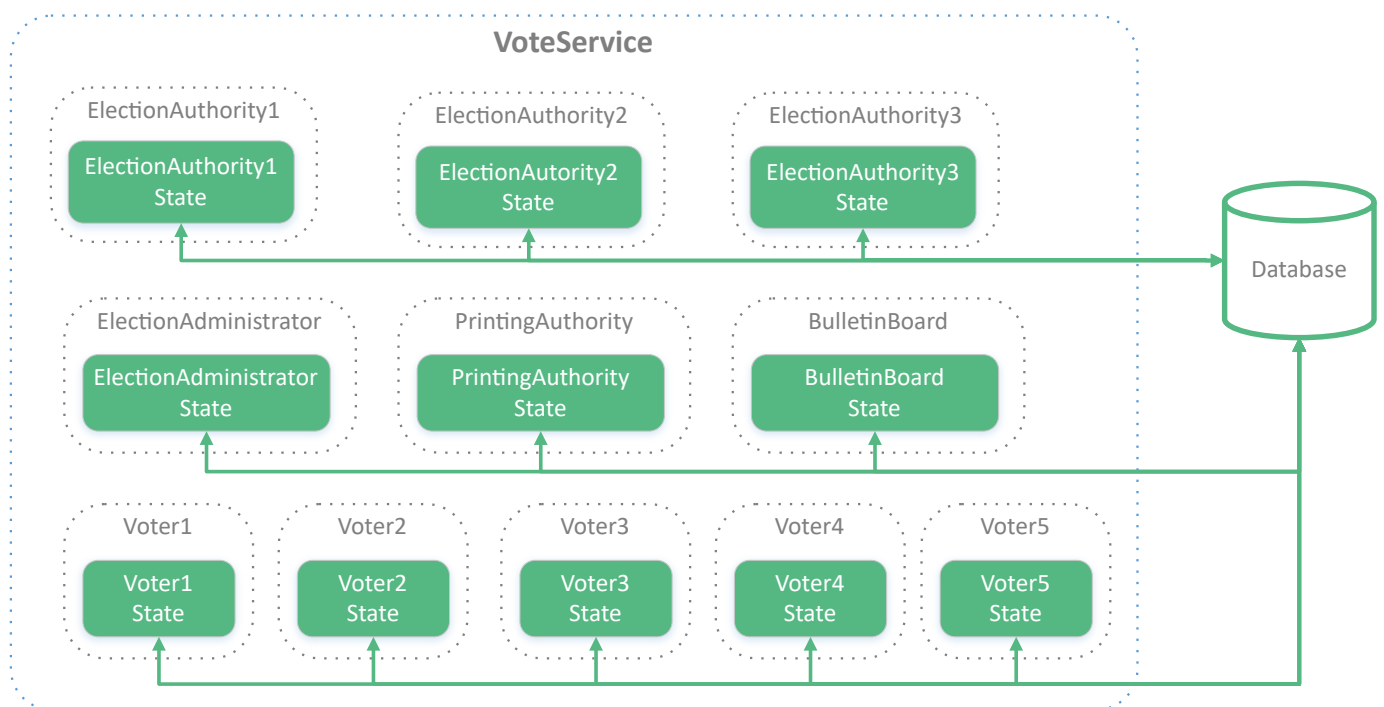


Figure 5.3: The internals of our VoteService: The functionality is divided into actor- and state-classes. The actors provide the functionality as described in the specification and by utilizing our crypto-library. The state is contained within the state classes that allow easy serialization both for persisting them to the database as well as to a JSON representation for the data synchronization.

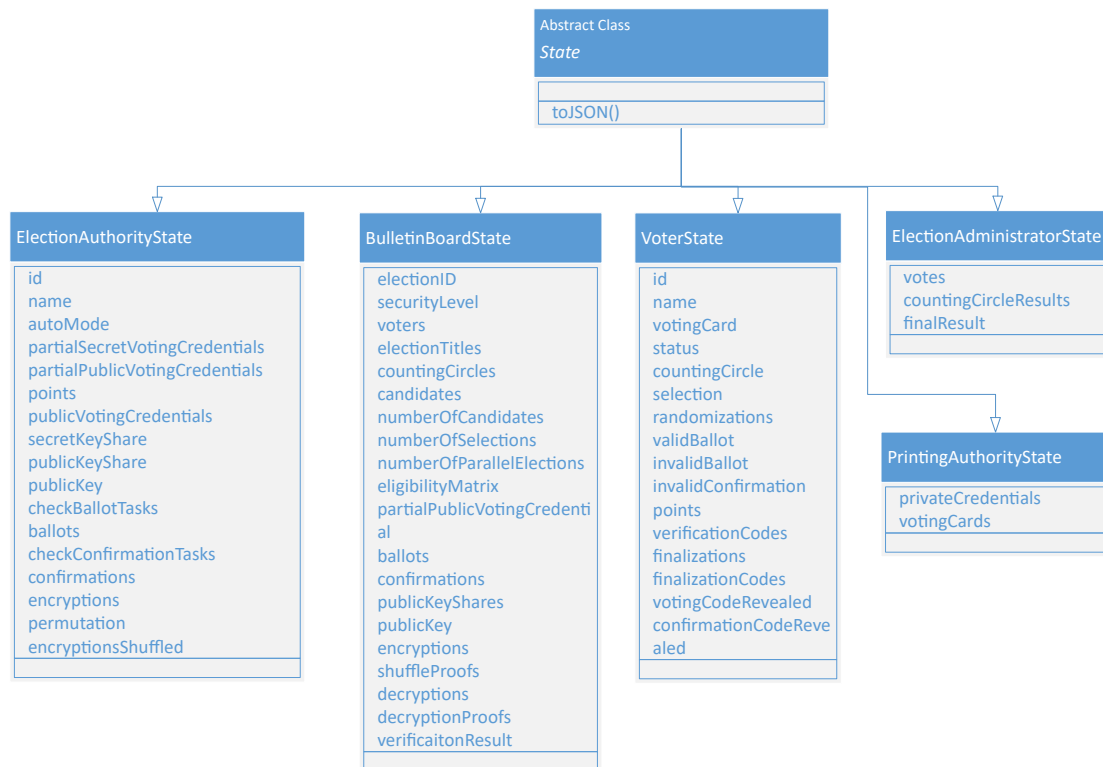


Figure 5.4: State classes

- **BulletinBoardState:** Holds all data that is publicly available on the bulletin board (the number of candidates, the tallied result)
- **ElectionAuthorityState:** Holds all data that an election authority knows (e.g. the list of ballots, the secret key of an election authority)
- **VoterState:** Since there is no distinction made between a voter and a voting client in our application, the VoterState contains the data of both the voter (e.g. the voting card) and the data typically known to the voting client (e.g. the points returned by the oblivious transfer)
- **PrintingAuthorityState:** Holds the data known to the printing authority (e.g. the list of all voters private credentials and the voting cards)
- **ElectionAdministratorState:** Holds all data known to the election administrator

Since our application supports that multiple users work on different election events concurrently, the state of an election event cannot be kept in volatile memory, but needs to be persisted between every single request. For this reason we evaluated different database systems and concepts. We decided against a relational database system that requires us to define a database schema as we wanted our state objects to be the only place where the schema is defined. This "code-first" approach makes it easier to apply changes to the protocol in future.

For our purpose, mongoDB seemed like a good choice. Since we do not need the ability to access and filter our data by arbitrary queries, but only need to be able to save and load

a state object of a particular election, we simply store the whole state as a binary string in a mongoDB collection. The only additional attribute that is saved to the database alongside with the serialized state is the electionId which denotes which election event a particular state belongs to. An election contains multiple VoterStates and ElectionAuthorityStates. Therefore, these two states additionally require an electionAuthorityId and a voterId.

```

1  _id: ObjectId("5a040ba19a7c4c40c8b310a7")
2  election : "5a040ba19a7c4c40c8b310a5"
3  state : Binary('gANjYXBwLm1vZGVscyslbGVjdGlvbG91dGVhcn0eVN0YXRlcKVsZW90aw9uQXV0aG9yaXR5U3RhdGUKcQApGXEbfXECKFgXAAAAA...')
4  authorityID : 0

```

Figure 5.5: Example: ElectionAuthority document collection

We described how the state classes are used to divide the data of the VoteService into smaller units. Similarly, the functionality of the VoteService is separated into classes, one for every actor of the protocol.

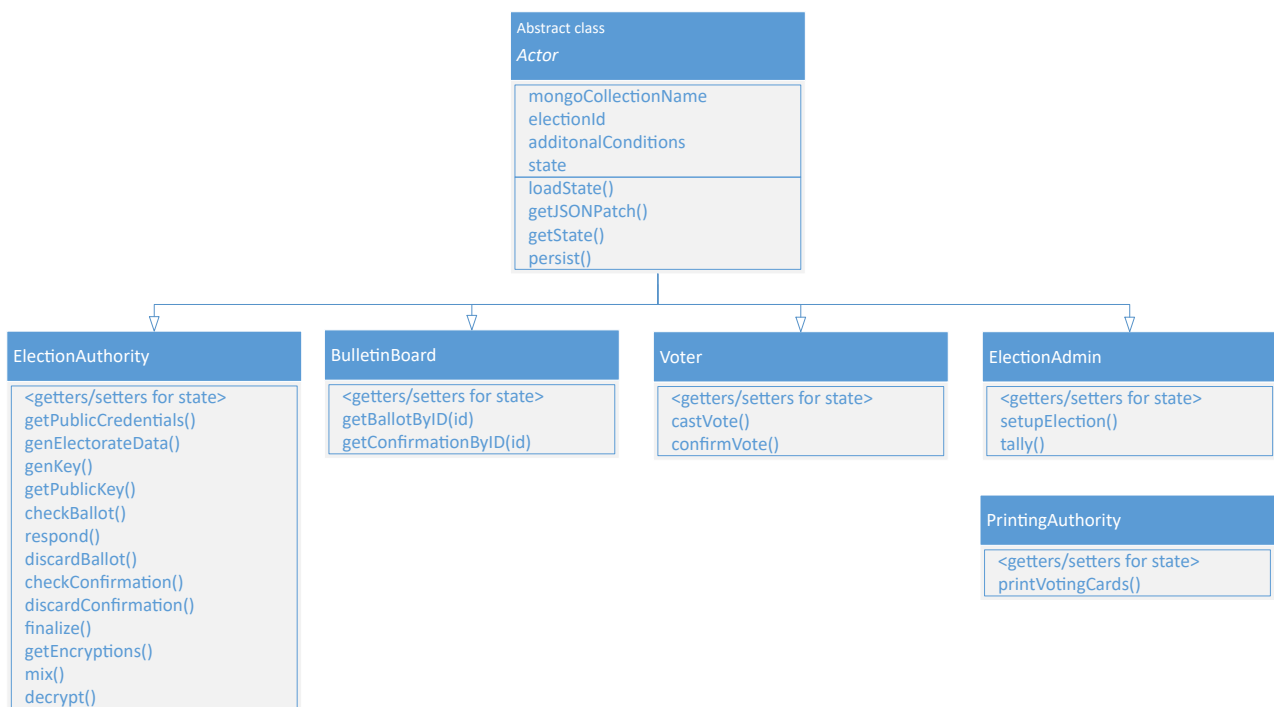


Figure 5.6: Actor classes

The common functionality, namely, a function for loading the corresponding states from the database and one for persisting the states to the database, are contained in an abstract base class.

5.3.2 Data-Sync Service

One of the big challenges of our application has been the synchronization of an election events state from the back-end to the clients local data store. As mentioned earlier, the web-applications contain a local datastore, which is structured the same way as the states of the VoteService. As per our requirements, we wanted to achieve real time data synchronization such that every web client observing a particular election, is informed of any change of this election events state. For this purpose we have used web-sockets which - opposed to the HTTP protocol - makes it possible to inform a web-client without having to rely on polling.

For the data synchronization we had to keep an eye on the performance of the data transfers since some state objects, especially the bulletin board and the election authority states grow big in size when they contain many ballots. We observed that the size of the whole state of an election event with 6 candidates and 10 voters, of which each has submitted at least one ballot and a confirmation can easily reach 600 kilobytes already. Admittedly, we haven't noticed any performance issues even with rather large election events, however, transferring the whole state of an election after every single mutation didn't seem like a proper solution.

Nevertheless, when a client connects to the data sync service for the first time, it needs to fetch the JSON representation of every state object of the VoteService. For this purpose we have implemented a "FullSync" method which will populate the clients local data store with the full state of an election.

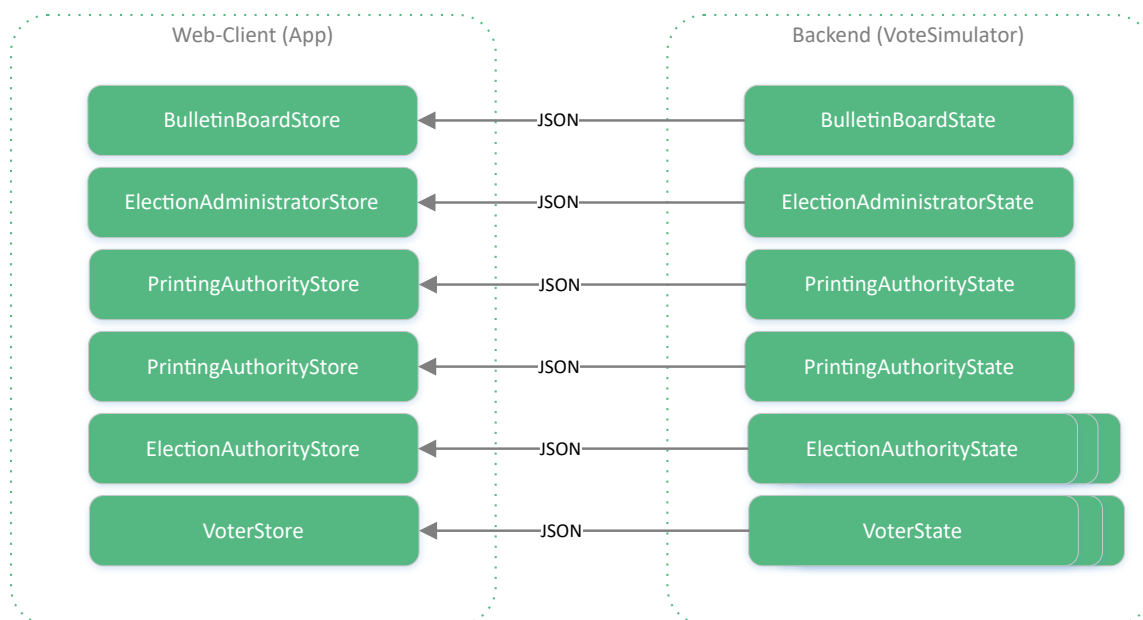


Figure 5.7: Datastores

After a client has populated his local data store, future manipulations on the backend are synchronized using the so called JSON Patch operations, which only contain the delta between the previous and the current state.

JSON Patch is a structure for describing how a JSON object can be modified / "patched".

The procedure is standardized and described in the RFC 6902 of the Internet Engineering Task Force (IETF). There exist JSON Patch implementations for many languages, including Python and JavaScript. We used JSON Patch to realize our incremental data synchronization as follows:

When the VoteService loads the state of its actor objects, it sets the `originalState` property of the actor to a deep copy of his state object. Mutations are always done only to the `state` property. Before calling the `persist()` method on an actor object, we use the Python JSON Patch library to determine the differences between the state and the `originalState`, to find out if and what variables have changed within the states. Additionally, we generate a set of JSON Patch operations that describes how the `originalState` could be patched, such that it becomes identical to the manipulated `state` object, by calling `make_patch(json.loads(self.originalState.toJSON())`

The result is an array of operations in JSON format that contains:

- The path of the manipulation
- The type of operation (replace, add, remove, ...)
- The new value (if required)

As an example: After casting a ballot, we might receive the following JSON patch:

```
▼ {election_authority_0: Array(1), printing_authority: Array(0),  
  ▼ bulletin_board: Array(1)  
    ► 0: {path: "/ballots/0", op: "add", value: {...}}  
  ► election_administrator: []  
  ▼ election_authority_0: Array(1)  
    ► 0: {path: "/checkBallotTasks/0", op: "add", value: {...}}  
  ► election_authority_1: []  
  ► election_authority_2: []  
  ► printing_authority: []  
  ► voters: [{...}]
```

Figure 5.8: JSON Patch example

These JSON patches are pushed to all the clients that need to receive the mutations and are applied to the local data-store which (under normal circumstances) contains the original state. After applying the JSON patch, the data store of all clients contains the same state of an election event as the back-end.

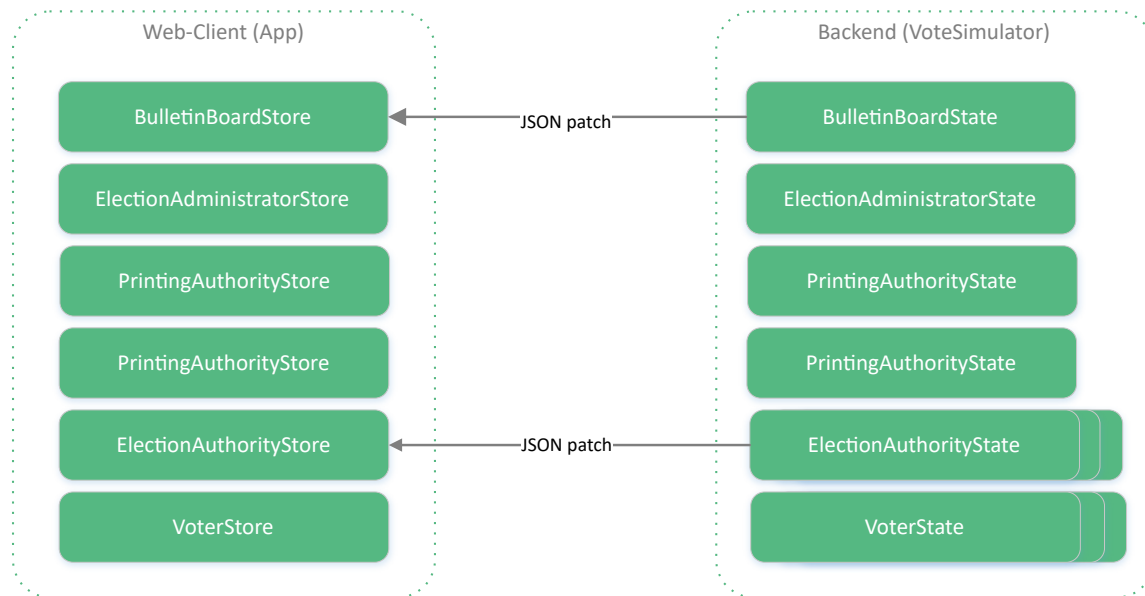


Figure 5.9: Data-Sync with JSON Patches

If for some reason, a web-client doesn't receive a JSON patch, its local data-store will no longer be corresponding to the actual data-store of the server. This might happen because of network issues and an interrupted websocket connection or if applying the JSON patch operations failed.

For this reason we have invented a data-store revision-number for every election event, which is incremented whenever some state is manipulated and persisted on the back-end. This revision number is sent alongside the JSON patches to the clients, and is also stored on the client side. Before applying the JSON patches, the client checks if his local data-stores revision number is exactly 1 version behind the servers data store revision, which normally will be the case. If the client detects that his revision number is 2 or more versions behind the servers, it will request a full-data synchronization over the DataSyncService to avoid out-dated, corrupted local data stores.

During development, we ran into an issue with the python JSON patch library "`python-json-patch v1.16`" that we have been using. During some special cases the generation of JSON patches failed and resulted in an exception when comparing objects where mutations have been made to arrays within dictionaries - a combination which often occurs in our data structures. After hours of debugging and analyzing the issue, we have figured out a temporary workaround and reported the issue ² to the repository of the developers of this library. A few days later a new version v1.20 of the library has been released which fixed our issue.

5.3.3 REST API

The third component of the backend is the REST API. Its responsibilities are to provide all the functionality of the VoteService to the webclients and trigger the data synchronization of

²<https://github.com/stefankoegl/python-json-patch/issues/74>

the DataSyncService. Every function required by the front-end, such as `castVote()` has a corresponding endpoint in the REST API service.

```
@main.route('/castVote', methods=['POST'])
@cross_origin(origin='*')
def castVote():
    data = request.json
    electionId = data["election"]
    selection = data["selection"]
    voterId = data["voterId"]
    votingCode = data["votingCode"]

    if len(selection) == 0:
        return make_error(400, "Empty selection")

    try:
        svc = VoteService(electionId) # prepare VoteService

        svc.castVote(voterId, selection, votingCode) # perform vote casting

        patches = svc.persist() # persist modified state and retrieve JSON patches

        syncPatches(electionId, SyncType.ROOM, patches) # send the JSON patches to all
        ↪ clients

    except Exception as ex:
        return make_error(500, str(ex))

    return json.dumps({'result': 'success'})
```

The API can be reached by sending a HTTP(S) POST request to our webserver hosting the backend services. The URL defines what function will be executed. For example: A POST request to `https://<server>:5000/castVote/` would call the above function. The required parameters are provided in the POST body.

As a first step, parameters are extracted from the POST request and validated if necessary. As the next step, a `VoteService` object is instantiated by passing the `electionId` to the constructor. The `VoteService` will internally load the states of the corresponding election event from the database and instantiate the actor objects such as the election authorities.

Now the `VoteService` can execute the function which the user intended to call, for example "CastVote". By calling the function `persist()`, the new state is written to the database and the JSON Patches of all mutations that "CastVote" has been made are determined, returned and can be sent to all clients by the help of the `DataSyncService`.

The following sequence diagram shows how the vote casting use case is implemented within the back-end and how all the components work together.

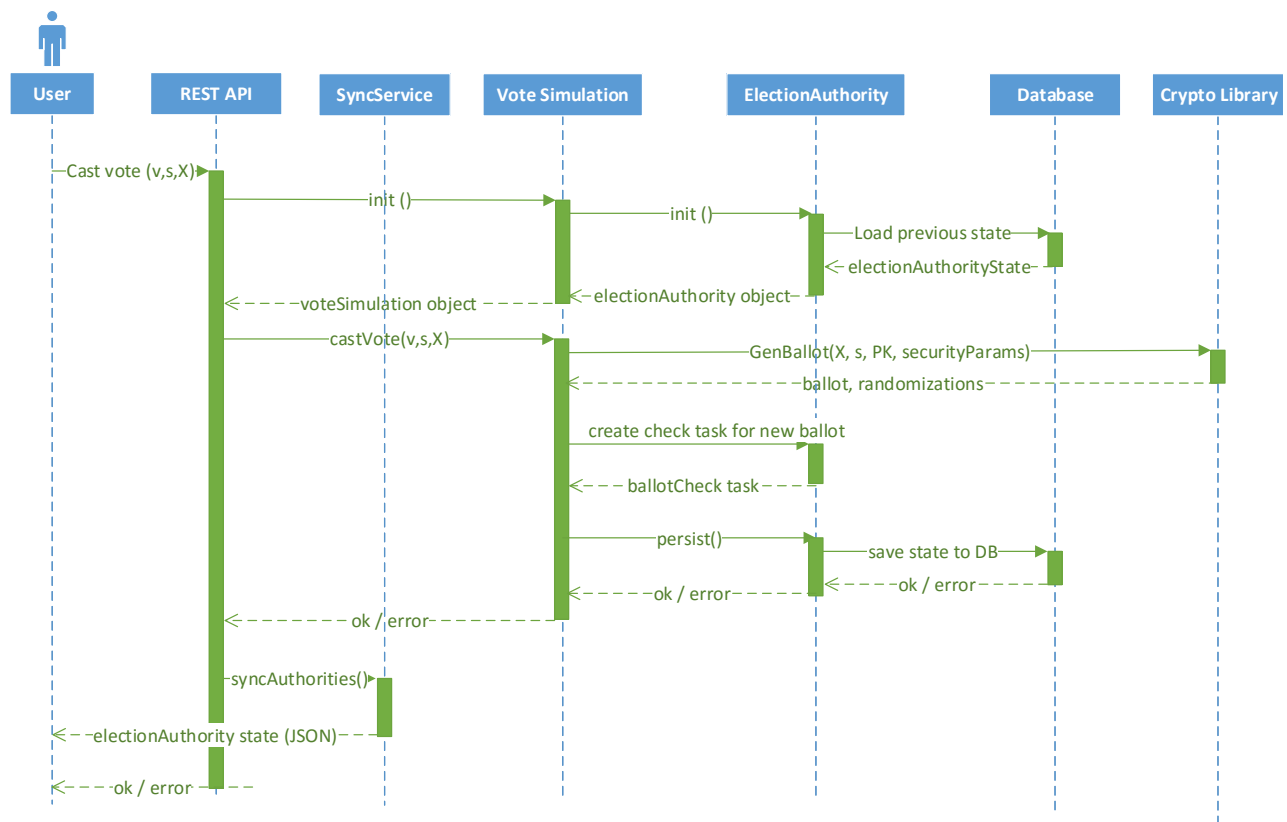


Figure 5.10: Vote casting sequence diagram

5.4 Crypto-Library

5.4.1 File structure

We decided to put every algorithm of the specification in its own file together with related unit tests. The files are structured according to the actors of the protocol, for example:

- **Common:** contains common cryptographic algorithms and the security parameters used by multiple algorithms
- **ElectionAuthority:** contains all the algorithms used by the election authority
- **PrintingAuthority:** contains all the algorithms used by the printing authority
- **VotingClient:** contains all the algorithms used by the voting client

- **ElectionAdministration**: contains all the algorithms used by the election administrator
- **Utils**: contains helper classes and miscellaneous utility functions

5.4.2 Public Parameters

There exist two types of public parameters:

The **security relevant parameters**, e.g:

- The order of the prime groups: $p, \ell p, \hat{p}$
- The length of the voting, confirmation, return and finalization codes
- The number of authorities: s

and **public election parameters**, e.g.:

- The size of the electorate: N_E
- The number of candidates: n
- The list of candidate descriptions: c

The security parameters are typically used within the algorithms and remain unchanged for a longer time period, whereas the public election parameters are different for every election event.

The object **SecurityParams** holds all security relevant parameters and is injected as an additional function argument to all algorithms. Several different **SecurityParams** objects are created initially, which contain all the parameters according to the recommendations in the CHVote specification document ("level 0" for testing purposes and "level 1" through "level 3" for productive use). For simple unit and debugging purposes, we can inject the "level 0" object while in production, level 1 - 3 are used.

The public election parameters on the other hand are directly passed to the algorithms by the calling party. If an algorithm needs to know certain election parameters (like the size of the electorate N_E), these values are typically derived from vectors that they have access to, so they do not require specific knowledge of these parameters.

5.4.3 Coding Style

The following source code sample shows a typical implementation of an algorithm (in this example, algorithm 7.18 according to the CHVote specification).


```

import gmpy2
from Utils.Utils import *
from Crypto.SecurityParams import *
from VotingClient.GetSelectedPrimes import GetSelectedPrimes
from VotingClient.GenQuery import GenQuery
from VotingClient.GenBallotProof import GenBallotProof
from Types import Ballot

def GenBallot(X_bold, s, pk, secparams):
    """
    Algorithm 7.18: Generates a ballot based on the selection s and the voting code X.
    ↪ The
    ballot includes an OT query a and a proof pi. The algorithm also returns the random
    values used to generate the OT query. These random values are required in Alg. 7.27
    to derive the transferred messages from the OT response, generated by Alg. 7.25.

    Args:
        X_bold (str): Voting Code  $X \in A_X \setminus X$ 
        s (list of int): Selection  $s = (s_1, \dots, s_k)$ 
        pk (mpz): ElGamal key  $pk \in G_p \setminus \{1\}$ 
        secparams (SecurityParams): Collection of public security parameters

    Returns:
        tuple:  $\alpha = (r, \text{Ballot}) = (r, (x_{\text{hat}}, a, b, \text{pi}))$ 
    """
    AssertMpz(pk)
    AssertList(s)
    AssertClass(secparams, SecurityParams)

    x = mpz(StringToInteger(X_bold, secparams.A_X))
    x_hat = gmpy2.powmod(secparams.g_hat, x, secparams.p_hat)

    q_bold = GetSelectedPrimes(s, secparams) # q = (q_1, ..., q_k)
    m = mpz(1)

    for i in range(len(q_bold)):
        m = m * q_bold[i]

    if m >= secparams.p:
        return None

    (a_bold, r_bold) = GenQuery(q_bold, pk, secparams)
    a = mpz(1)
    r = mpz(0)

    for i in range(len(a_bold)):
        a = (a * a_bold[i]) % secparams.p
        r = (r + r_bold[i]) % secparams.q

    b = gmpy2.powmod(secparams.g, r, secparams.p)
    pi = GenBallotProof(x, m, r, x_hat, a, b, pk, secparams)
    alpha = Ballot(x_hat, a_bold, b, pi)

    return (alpha, r_bold)

class GenBallotTest(unittest.TestCase):
    def testGenBallot(self):
        selection = [1,4] # select candidates with indices 1,4
        (ballot, r) = GenBallot(unittestparams.X, selection, unittestparams.pk,
        ↪ secparams_10)
        print(ballot)
        print(r)

if __name__ == '__main__':

```

```
unittest.main()
```

All algorithms contain a short description, which was taken as-is from the specification document, as well as a comment (Google-style documentation string), which can be used to automatically generate code documentation. The algorithm itself is implemented as close to the specification as possible, using the same variable names and (as far as the language supports it) similar control structures:

- The suffix `_bold` for emphasized (bold) variables, e.g. `p_bold` for \mathbf{p}
- The suffix `_hat` for variables with a hat, e.g. `a_hat` for \hat{a}
- The suffix `_prime` for variables with a prime, e.g. `a_prime` for a'
- etc.

Each file also contains unit test corresponding to the specific algorithm (if unit testing was considered useful for the particular algorithm).

The following example shows the similarities between the algorithm pseudo code and the actual implementation in Python:

Algorithm: `GenBallot(X, s, pk)`

Input: Voting code $X \in A_X^{\ell_X}$

Selection $s = (s_1, \dots, s_k)$, $1 \leq s_1 < \dots < s_k$

Encryption key $pk \in \mathbb{G}_q \setminus \{1\}$

$x \leftarrow \text{ToInteger}(X)$

$\hat{x} \leftarrow \hat{g}^x \bmod \hat{p}$

$\mathbf{q} \leftarrow \text{GetSelectedPrimes}(s)$

$m \leftarrow \prod_{i=1}^k q_i$

if $m \geq p$ **then**

return \perp

$(\mathbf{a}, \mathbf{r}) \leftarrow \text{GenQuery}(\mathbf{q}, pk)$

$a \leftarrow \prod_{i=1}^k a_i \bmod p$

$r \leftarrow \sum_{i=1}^k r_i \bmod q$

$b \leftarrow g^r \bmod p$

$\pi \leftarrow \text{GenBallotProof}(x, m, r, \hat{x}, a, b, pk)$

$\alpha \leftarrow (\hat{x}, \mathbf{a}, b, \pi)$

return (α, \mathbf{r})

```
x = mpz(StringToInteger(X_bold, secparams.A_X))
x_hat = gmpy2.powmod(secparams.g_hat, x, secparams.p_hat)
q_bold = GetSelectedPrimes(s, secparams)

m = mpz(1)
for i in range(len(q_bold)):
    m = m * q_bold[i]

if m >= secparams.p:
    return None

(a_bold, r_bold) = GenQuery(q_bold, pk, secparams)
a = mpz(1)
r = mpz(0)

for i in range(len(a_bold)):
    a = (a * a_bold[i]) % secparams.p
    r = (r + r_bold[i]) % secparams.q

b = gmpy2.powmod(secparams.g, r, secparams.p)
pi = GenBallotProof(x, m, r, x_hat, a, b, pk, secparams)
alpha = Ballot(x_hat, a_bold, b, pi)

return (alpha, r_bold)
```

5.4.4 Return Types

In most cases, when an algorithm returns more than a scalar datatype, tuples are used. Tuples allow to return multiple values from a function:

```
def foo():
    return (1, 2)
```

```
def main():  
    a, b = foo()
```

This way large parts of the source code looked very similar to the pseudo code in the CHVote specification. For more complex data types or return values that are used more than once, named tuples were used. The data type "namedtuple" is like a lightweight class and allows access to named properties.

```
Ballot = namedtuple("Ballot", "x_hat, a_bold, b, pi")  
  
def main():  
    Ballot b = getBallot()  
    x_hat = b.x_hat
```

By following this approach we could avoid having lots of container classes only used to pass data structures between the algorithms.

5.5 Front-End

The front-end was the most important component for our project, since we put focus mostly on the visualization and less on the actual e-voting system. Displaying the rather large amount of voting specific data and large numbers required a clean and well structured layout and a modular component design. Luckily, the framework we had chosen, VueJS, did very well in supporting us to meet exactly those requirements. We tried to follow the design patterns and best practices proposed by the makers of VueJS wherever possible.

In this section we will explain what concepts of VueJS we used and how we adapted them to our needs.

5.5.1 Components

Components are the basic building blocks of the VueJS framework. The application itself is a component, every page of the application is a component and the pages typically contain lots of components, one for every control like form controls, buttons or custom controls such as the ballot-list etc. The concept of components encourages to create reusable modules, provides a way to structure the application into smaller units and makes the resulting HTML template more expressive and easier to read.

We have created our own VueJS components for every control that we used more than once. For example the ballot list that is shown both in the bulletin board view as well as in the election authority view, the labels for displaying truncated large numbers or the cards used as our standard mean for displaying data have all been turned into custom components. One of the

beautiful features of VueJS components is the concept of slots. By defining one or multiple slots within a components template markup, it becomes possible from the parent of a component, to embed content into different locations (slots) within the components HTML template.

We have been using slots to create our card component, which can display information either as the main content of the card, or within an expandable area on the bottom of the card:

```
<DataCard title="Foo" :expandable=true>
  Just some text
  <p slot="expandContent">More complex content <BigIntLabel
    ↪ :mpzValue="publicKey"></BigIntLabel>
  </p>
</DataCard>
```

The first line "Just some text", which gets inserted into the default unnamed slot, could as well be passed as a string parameter to the data card component. However, as things are getting more complicated, one might like to place arbitrary HTML or even another VueJS component inside the expandable content of our datacard. In such cases, component parameters won't work as they only accept primitive data types. Slots on the other hand allow arbitrary content to be injected. The following code shows how the data-card component is implemented:

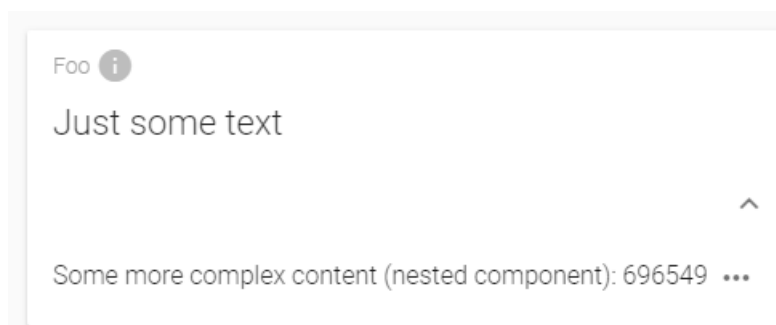


Figure 5.11: Vote casting sequence diagram

```
<template>
  <v-card class="dataCard">
    <ConfidentialityChip v-if="showConfidentiality" :type="confidentiality"
      ↪ class="confidentialityChip" />
    <v-card-title primary-title class="dataCardTitle">
      <div><span class="label grey--text">{{title}}
        <v-tooltip top>
          <v-icon v-if="!disableTooltip" color="grey lighten-1"
            ↪ slot="activator">info</v-icon><span>Programmatic tooltip</span>
        </v-tooltip></span>
      </div>
    </v-card-title>
    <v-card-text class="dataCardContent">
      <slot></slot>
    </v-card-text>
    <v-card-actions v-show="expandable">
      <v-btn icon @click.native="showExpander = !showExpander">
        <v-icon>{{ !showExpander ? 'keyboard_arrow_down' : 'keyboard_arrow_up'
          ↪ }}</v-icon>
      </v-btn>
    </v-card-actions>
  </v-card>
</template>
```

```
        </v-card-actions>
        <v-slide-y-transition v-show="expandable">
          <v-card-text v-show="showExpander">
            <slot name="expandContent">
            </slot>
          </v-card-text>
        </v-slide-y-transition>
      </v-card>
    </template>
    <script>
      import { mapState } from 'vuex'

      export default {
        data: function () {
          return {
            showExpander: false
          }
        },
        computed: {
          ...mapState({
            showConfidentiality: state => state.showConfidentiality
          })
        },
        props: {
          title: {
            type: String,
            required: true,
            default: 'Title'
          },
          expandable: {
            type: Boolean,
            required: true,
            default: false
          },
          confidentiality: {
            type: String,
            required: true,
            default: 'public'
          },
          disableTooltip: {
            type: Boolean,
            required: false,
            default: false
          }
        },
        mounted () {
        }
      }
    </script>
```

The first part within the `<template>` tag describes the HTML markup as well as the slots that we have just mentioned.

The `<script>` tag contains the actual logic of the component, similarly to the controller in other SPA frameworks. The `data` object contains variables which are defined and valid only locally within the component. The `computed` object maps variables from our central data-store to a local variable which is reactively bound to the data-store. Whenever the value of the given variable changes in the data-store, the computed property is automatically updated. From the template, we can access both local data as well as computed properties. Computed properties

can also be used whenever a local variable needs to be formatted, or in some way manipulated before it is displayed in the components template.

The third source of data are properties ('props'), which are passed as arguments from the parent component. They are typically used to defined options for a component.

Additionally, components may contain **methods**, typically used for event-handlers like button clicks and event hooks like **mounted**, **beforeDestroy**, **beforeCreated** to influence the components creation/destruction at different times during a components lifecycle.

5.5.2 Centralized Data-Store & Flux Pattern

One of the big challenges regarding the architecture of our front-end, was about how and where we would store all the data of an election. Clearly, since we already had divided an election events data into one state for every actor and given that every actor also has it's corresponding view in our front-end, simply storing the data to the respective component has been our first thought. Although a voter mainly needs to access his own data contained in the voters state, some data must be shared between multiple components, for example the information on the bulletin board.

Since we wanted to avoid having to keep data redundant in multiple components, we have decided to use the Flux design pattern in our front-end. The basic idea of the flux pattern is to have a single, central data source where all the data is stored and which all components have access to. This single data source is called a "store" by Flux terminology. VueJS has it's own implementation of the Flux pattern called Vuex. Another important concept is that components can freely access the data in the store, however, they are not allowed to change data, at least not directly. Instead, if a component wants to manipulate data in the store, this has to be done by calling **mutation functions**. Not allowing direct manipulation of the store makes it much easier to keep track of where mutations came from.

Our web applications data store is divided into multiple modules, one data store module for each corresponding state of the back-end. In reality, all these data stores are part of one single data store, but having multiple modules allows us to structure the mutation and getter functions and help to avoid naming conflicts by having separate namespaces for every module.

The data-store can be accessed from any component by defining a computed property. If the computed properties have to do perform some formatting, aggregation, filtering etc. on the state variables and are used from within multiple components, it is also possible and recommended to write getter-functions in the data-store to avoid redundancy.

5.5.3 Internationalization (i18n)

TODO: I18n

5.5.4 Development Environment

TODO: Webpack, ESLint etc. beschreiben

5.6 Challenges

In this section we describe some of the challenges that we encountered while developing our application.

5.6.1 Websocket subscription concept

Since our application allows that multiple elections exist at the same time, the question has arisen how we could handle that only those clients who are actually observing a particular election receive web-socket messages when an action has been triggered in the respective election.

We have seen similarities between our problem and the one a typical chat has that consists of multiple chat-rooms in which only those users should be notified of new posts that have actually joined the room. We have adapted this "chat-room" concept to our problem by defining a room to be equal to an election.

We can assume that all the pages that actually show data of an election require the election id to be passed as part of the URL. For example: `/BulletinBoard/1` is the route to reach the BulletinBoard view of the election with id 1.

Whenever a route is called that contains the argument 'electionId', we need to make sure that the client has joined this election's room. We therefore set a global variable called

6 Conclusion

During the first few weeks we felt as if we have been thrown into cold water. Reading and understanding the protocol wasn't easy at first, because we had to get used to the notation and memorize a large amount of variables used by the many algorithms. While some of the cryptographic primitives were taught in previous courses, most of them were new and unknown to us. We focused on getting a good understanding of the protocol on a higher level rather than learning about each and every algorithm in detail, as this was sufficient for implementing and understanding the protocol.

Additionally, programming algorithms isn't something we are doing on a daily basis. Therefore, the first few algorithms took us quite some time to implement. After a few weeks, we could greatly increase our productivity and in the end, we could implement even the larger algorithms in not much more time than the simple ones in the beginning of the project.

From our perspective, the project has been extremely interesting and we are still impressed by the ideas presented and specified in the CHVote specification. From simply implementing the protocol we could learn a lot about the CHVote protocol and E-Voting in general and could improve both our knowledge of more advanced cryptographic topics and get practise in implementing cryptographic algorithms.

6.1 Python issues

During the project we have experienced a few issues with the programming language that we used to implement the specification in, Python. In particular, we have observed the following issues:

- Performance issues due to Python being an interpreted language
- Function overhead: function calls in Python seem to be quite slow
- Strongly dynamic typing vs. static typing: the Python interpreter needs to inspect every single object during run time (be it an integer or a more complex object)
- The *BigInteger* library surprisingly isn't as fast as using directly the GMP library
- Larger projects tend to turn out messy
- Little to no standard documentation regarding project structure

- No real standard for unit testing, documentation generation etc.

For detailed information regarding the performance issues that we have experienced see [3] and [4]. Based on the reasons above we would not recommend to use Python for the use in similar or larger project. Python is indeed a very handy language to write quick prototypes and proof of concepts, but issues become more frequent in larger projects.

6.2 Reflection

Bibliography

- [1] "CHVote System Specification", by Rolf Haenni, Reto E. Koenig, Philipp Locher and Eric Dubuis, April 11, 2017.
- [2] "'Anforderungskatalog für eidgenössische Volksabstimmungen mit der elektronischen Stimmabgabe'", by Bundeskanzlei BK, June 07.2014)
- [3] "Why Python is Slow: Looking Under the Hood", by Jake VanderPlas, see <http://jakevdp.github.io/blog/2014/05/09/why-python-is-slow/>
- [4] "Python speed: performance tips", from the official Python wiki, see <https://wiki.python.org/moin/PythonSpeed/PerformanceTips>

7 Appendix

7.1 Use Cases

Table 7.1: Use Case «Create new election»

Use Case	Create new election
Primary Actor	User
Description	The system allows to create new elections
Precondition	The system shows the available elections in a list
Main path (M)	<ol style="list-style-type: none"> 1. User clicks on "create election" 2. System demands a name for the election 3. User is redirected to the election overview page

Table 7.2: Use Case «Set-up an election event»

Use Case	Set-up an election event
Primary Actor	Election Administrator
Description	The election administrator can set up the election. This involves the generation of the cryptographic electorate data in the backend
Precondition	An new election has been created
Postcondition	The election has the status "Printing"
Main path (M)	<ol style="list-style-type: none"> 1. Election Administrator visits the "Election Administrator"-view of a new election. 2. The system demands the following information: <ul style="list-style-type: none"> • Number of parallel elections • Candidates per election • Number of possible selections per election event • Number of voters • Counting circles of the voters 3. User clicks on "Generate"

Table 7.3: Use Case «Printing of voting cards»

Use Case	Printing of voting cards
Primary Actor	Printing Authority
Description	The printing authority generates voting cards
Precondition	The election has the status "‘Printing”’
Postcondition	The election has the status "‘Delivery”’
Main path (M)	<ol style="list-style-type: none"> 1. The election administrator visits the "‘Printing Authority”’-view of an election. 2. The election administrator clicks on "‘Print Voting Cards”’ 3. A list of all voters is displayed 4. The election administrator can select a voter to see his voting card

Table 7.4: Use Case «Delivery of voting cards»

Use Case	Delivery of voting cards
Primary Actor	Printing Authority
Description	The printing authority can send the voting cards to the voters
Precondition	The election has the status "‘Delivery”’
Postcondition	The election has the status "‘Election Phase”’
Main path (M)	<ol style="list-style-type: none"> 1. The election administrator visits the "‘Printing Authority”’-view of an election. 2. The election administrator clicks on "‘Deliver Voting Cards”’ 3. Within the voters-view, the voting card shows up for every voter

Table 7.5: Use Case «Confirmation of a vote»

Use Case	Confirmation of a vote
Primary Actor	Voter
Description	The voter can confirm his vote by verifying the verification codes and entering his confirmation code
Precondition	<ul style="list-style-type: none"> • The election has the status "Election Phase" • A voter is selected in the "Voter"-view • The voter has the status "Confirmation Phase"
Postcondition	The first election authority receives a "Check-confirmation task"
Main path (M)	<ol style="list-style-type: none"> 1. The voter visits the "Voter"-view and select his voter object 2. The system displays the verification codes of the selected candidates 3. The voter must manually verify that the displayed codes match the verification codes of the selected candidates on his voting card 4. The system demands the confirmation code 5. The voter clicks on "Confirm vote"

Table 7.6: Use Case «Checking a ballot»

Use Case	Checking a ballot
Primary Actor	Election Authority
Description	The election authority can verify the validity of a ballot and respond to the voters query
Precondition	<ul style="list-style-type: none"> • The election has the status "Election Phase" • The currently selected election authority has a new "Check ballot task"
Postcondition	<ul style="list-style-type: none"> • The next election authority receives a "Check ballot task" • If this election authority was the last one, and the ballot was valid, the voter now has the status "Confirmation Phase"
Main path (M)	<ol style="list-style-type: none"> 1. The user visits the "Election Authority"-view and select one of the available election authorities that has new "Check ballot task" 2. The system displays the query, the ballot proof and the voting credential of the voter 3. The user click on "Check validity" 4. The system displays the result of the validity check 5. The user clicks on "Respond"

Table 7.7: Use Case «Checking a confirmation»

Use Case	Checking a confirmation
Primary Actor	Election Authority
Description	The election authority can verify the validity of a confirmation and respond to the voters query
Precondition	<ul style="list-style-type: none"> • The election has the status "Election Phase" • The currently selected election authority has a new "Check ballot task"
Postcondition	<ul style="list-style-type: none"> • The next election authority receives a "Check confirmation task" • If this election authority was the last one, and the confirmation was valid, the voter now has the status "Finalization Phase"
Main path (M)	<ol style="list-style-type: none"> 1. The user visits the "Election Authority"-view and select one of the available election authorities that has new "Check confirmation task" 2. The system displays information about the confirmation 3. The user click on "Check validity" 4. The system displays the result of the validity check 5. The user clicks on "Finalize"

Table 7.8: Use Case «Mixing»

Use Case	Mixing
Primary Actor	Election Authority
Description	Every election authority can perform the mixing on the extracted list of encryptions
Precondition	<ul style="list-style-type: none"> • The election has the status "Mixing" • The previous election authority has already performed the mixing
Postcondition	<ul style="list-style-type: none"> • The next election authority is able to mix
Main path (M)	<ol style="list-style-type: none"> 1. The user visits the "Election Authority"-view and select one of the available election authorities that hasn't mixed before 2. The system displays the list of encryptions of the previous election authority (or the first one in case the first election authority is selected) 3. The user clicks on "Mix" 4. The new, mixed list of encryptions is added to the known data of this election authority

Table 7.9: Use Case «Decryption»

Use Case	Decryption
Primary Actor	Election Authority
Description	Every election authority can perform the (partial) decryption
Precondition	<ul style="list-style-type: none"> • The election has the status "«Decryption»" • The previous election authority has already performed the decryption
Postcondition	<ul style="list-style-type: none"> • The next election authority is able to decrypt
Main path (M)	<ol style="list-style-type: none"> 1. The user visits the "«Election Authority»"-view and select one of the available election authorities that hasn't decrypted before 2. The system displays the list of encryptions 3. The user clicks on "«Decrypt»" 4. The list of partial decryptions is added to the known data of this election authority

Table 7.10: Use Case «Tallying»

Use Case	Tallying
Primary Actor	Election Administrator
Description	The election administrator can perform the tallying and view the final result
Precondition	The election has the status "«Tallying»"
Postcondition	The has the status "«Finished»"
Main path (M)	<ol style="list-style-type: none"> 1. The user visits the "«Election Administrator»"-view 2. The user clicks on "«Tally»" 3. The final result is added to the known data of the election administrator